

HW4

March 15, 2019

0.0.1 Homework 4


Erin Witmer CSC 440

[8.3] Given a decision tree, you have the option of (a) converting the decision tree to rules and then pruning the resulting rules, or (b) pruning the decision tree and then converting the pruned tree into rules. What advantage does (a) have over (b)? A decision tree can be converted to IF-THEN rules by tracing the path from the root node to each leaf in the tree. If the tree is pruned before the rules are created, as is suggested in case (b), then any path dependent on the trimmed node will also be trimmed. First creating the rules, and then trimming them as proposed in (a) allows for more precision when trimming. Leaves dependent on a specific preceding node may be trimmed, while other leaves dependent on that same preceding node may be kept. However, it is important to note the higher computational complexity associated with building a full tree, converting to rules and then trimming. Regardless, method (a) generally generates a more precise and accurate set of rules and is worth the computational cost.

[8.5] Given a 5-GB data set with 50 attributes (each containing 100 distinct values) and 512 MB of main memory in your laptop, outline an efficient method that constructs decision trees in such large data sets. Justify your answer by rough calculation of your main memory usage. The use of the RainForest algorithm would be most effective in this case. RainForest adapts to the amount of main memory available and applies to any decision tree induction algorithm. The first step is to construct an AVC-set for each attribute. If the total data set is 5MB, on average each attribute would be 100MB. The process would be to extract one attribute, calculate the AVC, and store that set in main memory. On average, if you assume you are storing the value/class count as an integer (4 bytes), there are 100 values per attribute, and assume there are 2 class outcomes, an individual AVC set would require approximate $(100 \times 2 \times 4 \text{ bytes}) = 800 \text{ bytes}$. So the data holding the aggregate information from which the decision tree can be constructed is only 40 KB. Since this is the aggregate data, every level of the tree will be some derivative of this data. Utilizing the basic algorithm for inducing a decision tree, you could loop through many times without running out of memory. If the tree got very deep and space was an issue, tree pruning could take place while the tree is being generated and only summary information could be retained at each level.

[8.7] The following table consists of training data from an employee database. The data have been generalized....

```
In [1]: data = {'t0': ['sales', 'senior', '3135', '46K50K', 30],
               't1': ['sales', 'junior', '2630', '26K30K', 40],
               't2': ['sales', 'junior', '3135', '31K35K', 40],
```



```

't3': ['systems', 'junior', '2125', '46K50K', 20],
't4': ['systems', 'senior', '3135', '66K70K', 5],
't5': ['systems', 'junior', '2630', '46K50K', 3],
't6': ['systems', 'senior', '4145', '66K70K', 3],
't7': ['marketing', 'senior', '3640', '46K50K', 10],
't8': ['marketing', 'junior', '3135', '41K45K', 4],
't9': ['secretary', 'senior', '4650', '36K40K', 4],
't10': ['secretary', 'junior', '2630', '26K30K', 6]}

```

(a.) How would you modify the basic decision tree algorithm to take into consideration the count of each generalized data tuple? The relative weights of each of the tuples in the entropy calculation needs to be adjusted to reflect the proportion of the count to total employees (sum of count).

(b.) Use your algorithm to construct a decision tree from the given data.

```

In [2]: import pandas as pd
import math
from functools import reduce

def entropy(a,b):
    sum = a+b
    total = -(a/sum)*math.log2((a/sum))-(b/sum)*math.log2((b/sum))

```

```

        return total

df = pd.DataFrame.from_dict(data, orient='index', columns=['dept', 'class', 'age', 'sa
class_count = df.groupby('class').sum()
print(class_count)
head = entropy(113,52)

count
class
junior    113
senior     52

```

```

In [3]: # department
dept_count = df.groupby(['dept']).sum()
print(dept_count)

```

```

count
dept
marketing    14
sales        110
secretary     10
systems       31

```

```

In [4]: sales = 110/165
        systems = 31/165
        marketing = 14/165
        secretary = 10/165

dept_outcome = df.groupby(['dept', 'class']).sum()
print(dept_outcome)

```

```

count
dept  class
marketing junior    4
      senior    10
sales   junior    80
      senior    30
secretary junior    6
      senior    4
systems  junior    23
      senior    8

```

```

In [5]: dept_gain = head - (sales*entropy(80,30)
                          +systems*entropy(23,8)
                          +marketing*entropy(4,10)
                          +secretary*entropy(6,4))

print(dept_gain)

```

0.048606785991983426

```
In [6]: # age
        age_count = df.groupby(['age']).sum()
        print(age_count)
```

	count
age	
2125	20
2630	49
3135	79
3640	10
4145	3
4650	4

```
In [7]: early20=20/165
        late20=49/165
        early30=79/165
        late30=10/165
        early40=3/165
        late40=4/165

        age_outcome = df.groupby(['age','class']).sum()
        print(age_outcome)
```

		count
age	class	
2125	junior	20
2630	junior	49
3135	junior	44
	senior	35
3640	senior	10
4145	senior	3
4650	senior	4

```
In [8]: age_gain = head - (early30*entropy(44,35))
        print(age_gain)
```

0.4247351209783661

```
In [9]: # salary
        salary_count = df.groupby(['salary']).sum()
        print(salary_count)
```

	count
salary	

26K30K	46
31K35K	40
36K40K	4
41K45K	4
46K50K	63
66K70K	8

```
In [10]: high20=46/165
         low30=40/165
         high30=4/165
         low40=4/165
         high40=63/165
         high60=8/165
```

```
salary_outcome = df.groupby(['salary','class']).sum()
print(salary_outcome)
```

		count
salary	class	
26K30K	junior	46
31K35K	junior	40
36K40K	senior	4
41K45K	junior	4
46K50K	junior	23
	senior	40
66K70K	senior	8

```
In [11]: salary_gain = head - (high40*entropy(23,40))
         print(salary_gain)
```

0.5375181264158646

```
In [12]: # The first level of the decision tree is by salary
         salary_age = df.groupby(['salary','age','class']).sum()
         print(salary_age)
```

			count
salary	age	class	
26K30K	2630	junior	46
31K35K	3135	junior	40
36K40K	4650	senior	4
41K45K	3135	junior	4
46K50K	2125	junior	20
	2630	junior	3
	3135	senior	30
	3640	senior	10

```
66K70K 3135 senior      5
      4145 senior      3
```

```
In [13]: salary_dept = df.groupby(['salary', 'dept', 'class']).sum()
        print(salary_dept)
```

salary	dept	class	count
26K30K	sales	junior	40
	secretary	junior	6
31K35K	sales	junior	40
36K40K	secretary	senior	4
41K45K	marketing	junior	4
46K50K	marketing	senior	10
	sales	senior	30
	systems	junior	23
66K70K	systems	senior	8

The first level of the decision tree is a split by salary, because the information gain is the highest based on that split. There is only one ambiguous category after the salary split (46k...50k), and a further split based on either age (21...25 -> junior, 26...30 -> junior, 31...35 -> senior, 36...40 -> senior) or department (marketing -> senior, sales -> senior, systems -> junior) will result in a pure split of the data.

(c.) Given a data tuple having the values "systems," "26...30," and "46-50K" for the attributes department, age, and salary, respectively, what would a naive Bayesian classification of the status for the tuple be?

```
In [14]: p_senior = 52/165
        p_junior = 113/165

        p_senior_systems = 8/52
        p_junior_systems = 23/113

        p_senior_26 = 1/52 # modifier
        p_junior_26 = 49/113

        p_senior_46 = 40/52
        p_junior_46 = 23/113

        senior = p_senior * p_senior_systems * p_senior_26 * p_senior_46
        junior = p_junior * p_junior_systems * p_junior_26 * p_junior_46

        print(senior)
        print(junior)
```

```
0.0007172314864622557
0.012302997078625555
```

The Bayesian classification would be junior.

[8.12] The data tuples of Figure 8.25 are sorted by decreasing probability value, as returned by a classifier. For each tuple, compute the values for the number of true positives (TP), false positives (FP), true negatives (TN), and false negatives (FN). Compute the true positive rate (TPR) and false positive rate (FPR). Plot the ROC curve for the data.

```
In [15]: roc = pd.read_excel('roc.xlsx', index_col=0, dtype={'TPR': float, 'FPR': float})
        roc
```


```
Out[15]:
```

	Class	Prob.	TP	FP	TN	FN	TPR	FPR
Tuple								
1	P	0.95	1	0	5	4	0.2	0.0
2	N	0.85	1	1	4	4	0.2	0.2
3	P	0.78	2	1	4	3	0.4	0.2
4	P	0.66	3	1	4	2	0.6	0.2
5	N	0.60	3	2	3	2	0.6	0.4
6	P	0.55	4	2	3	1	0.8	0.4
7	N	0.53	4	3	2	1	0.8	0.6
8	N	0.52	4	4	1	1	0.8	0.8
9	N	0.51	4	5	0	1	0.8	1.0
10	P	0.40	5	5	0	0	1.0	1.0

```
In [16]: import matplotlib.pyplot as plt
        tpr = roc.iloc[:,6].values
        fpr = roc.iloc[:,7].values
```

```
In [17]: plt.plot(fpr, tpr, 'o', color='black');
        plt.xlabel('false positive rate')
        plt.ylabel('true positive rate')
```

```
Out[17]: Text(0,0.5,'true positive rate')
```



[8.14] Suppose that we want to select between two prediction models, M1 and M2. We have performed 10 rounds of 10-fold cross-validation on each model, where the same data partitioning in round i is used for both M1 and M2. The error rates obtained for M1 are 30.5, 32.2, 20.7, 20.6, 31.0, 41.0, 27.7, 26.0, 21.5, 26.0. The error rates for M2 are 22.4, 14.5, 22.4, 19.6, 20.7, 20.4, 22.1, 19.4, 16.2, 35.0. Comment on whether one model is significantly better than the other considering a significance level of 1%.

```
In [18]: import numpy as np
m1 = [30.5, 32.2, 20.7, 20.6, 31.0, 41.0, 27.7, 26.0, 21.5, 26.0]
m2 = [22.4, 14.5, 22.4, 19.6, 20.7, 20.4, 22.1, 19.4, 16.2, 35.0]
```

```
In [19]: def var_diff(m1,m2):
    total = 0
    m1_avg = np.average(m1)
    m2_avg = np.average(m2)
    k = len(m1)
    d = m1_avg - m2_avg

    for i in range(k):
        add = (m1[i]-m2[i]-(m1_avg-m2_avg)) ** 2
        total += add

    total = total / k
    return d/np.sqrt(total / k)
```



```
var_diff(m1,m2)
```

```
Out[19]: 2.4712371600876786
```

```
In [20]: from scipy.stats import t
```

```
# two sided t-test with k-1 dof  
t.interval(0.99, df=9)
```

```
Out[20]: (-3.2498355440153697, 3.2498355440153697)
```

The p-value (2.47) is not greater than 3.25 or less than -3.25, therefore we cannot reject the null hypothesis that one model is significantly better than the other at a significance level of 1%.